

KAGGLE COMPETITION

Thea Boge

Dataset

From 45 raw columns → 10 relevant, interpretable features

Objective: “condition”

Textual

- title
- descriptions

Contains seller language that directly expresses product condition (e.g. “nuevo”, “sin uso”, “usado”)

Numerical

- price
- base_price
- sold_quantity
- available quantity

Capture pricing patterns and product demand, which differ between new and used items.

Categorical


- listing_type_id
- buying_mode
- category_id

Describe how items are sold and categorized, adding contextual variation.

Text Cleaning and Preprocessing

Implemented two functions:

- clean_text()
 - lowercased text
 - removed punctuation and special characters
 - normalized accented characters
- preprocess()
 - combine titles and descriptions
 - apply the clean_text function
 - feature engineering for the numeric and categorical variables



Further preprocessing

Three feature groups for later preprocessing techniques.


Preprocessing techniques:

- TF-IDF Vectorizer
- StandardScaler
- OneHotEncoder

```
# Feature columns
text_col = "text"
num_cols = ["price", "base_price", "sold_quantity", "available_quantity",
            "price_ratio", "sold_ratio", "price_diff", "log_price"]
cat_cols = ["listing_type_id", "buying_mode", "category_id"]
```

Support Vector Machines


- Baseline model for comparison
- Linear SVM: fast, interpretable, effective for high-dimensional sparse data
- Accuracy: 0.8692
- Non-linear kernel (RBF) could model complex patterns, but too heavy for this dataset.



XGBoost

- Gradient boosting algorithm
- Strong performance on mixed features
- Histogram tree method: faster and memory-efficient
- Accuracy: 0.8828

Moving forward → improving the XGBoost



<https://www.geeksforgeeks.org/machine-learning/xgboost/>

Improvement

- ColumnTransformer
- Added Pipeline
- Tuned hyperparameters
 - learning rate
 - max_depth
 - n_estimators
- Added feature engineering to preprocess() to capture economic behaviour

```
# XGBoost model
xgb_model = XGBClassifier(
    n_estimators=7000,
    learning_rate=0.03,
    max_depth=9,
    subsample=0.85,
    colsample_bytree=0.7,
    reg_lambda=1.5,
    reg_alpha=0.5,
    tree_method="hist",
    random_state=42,
    n_jobs=-1,
    eval_metric="logloss"
)
```

Improved model → 0.89133



```
df["price_ratio"] = df["price"] / (df["base_price"] + 1)
```

```
df["log_price"] = np.log1p(df["price"].fillna(0))
df["log_sold"] = np.log1p(df["sold_quantity"].fillna(0))
```

```
df["sold_ratio"] = df["sold_quantity"] / (df["available_quantity"] + 1)
```

Final Model Validation

XGBoost with histogram-based tree method, combining TF-IDF text, scaled numeric, and one-hot categorical features



■ Limitations and further work

- GridSearchCV() or Optuna for model tuning
 - Automate the search for optimal hyperparameters instead of manual tuning.
- Test additional boosting frameworks (e.g., LightGBM, CatBoost)
 - Compare performance and training efficiency on mixed feature data.
- Use richer text embeddings
 - Replace TF-IDF with word embeddings such as FastText or BERT to capture semantic meaning.
- Feature importance and interpretability
 - Analyze the most influential features to better understand what drives the model's predictions.

Gracias por su atención!