

# KAGGLE COMPETITION

---

Thea Boge

# Introduction

1

**DATA UNDERSTANDING AND FEATURE SELECTION**

2

**BASELINE MODEL**

3

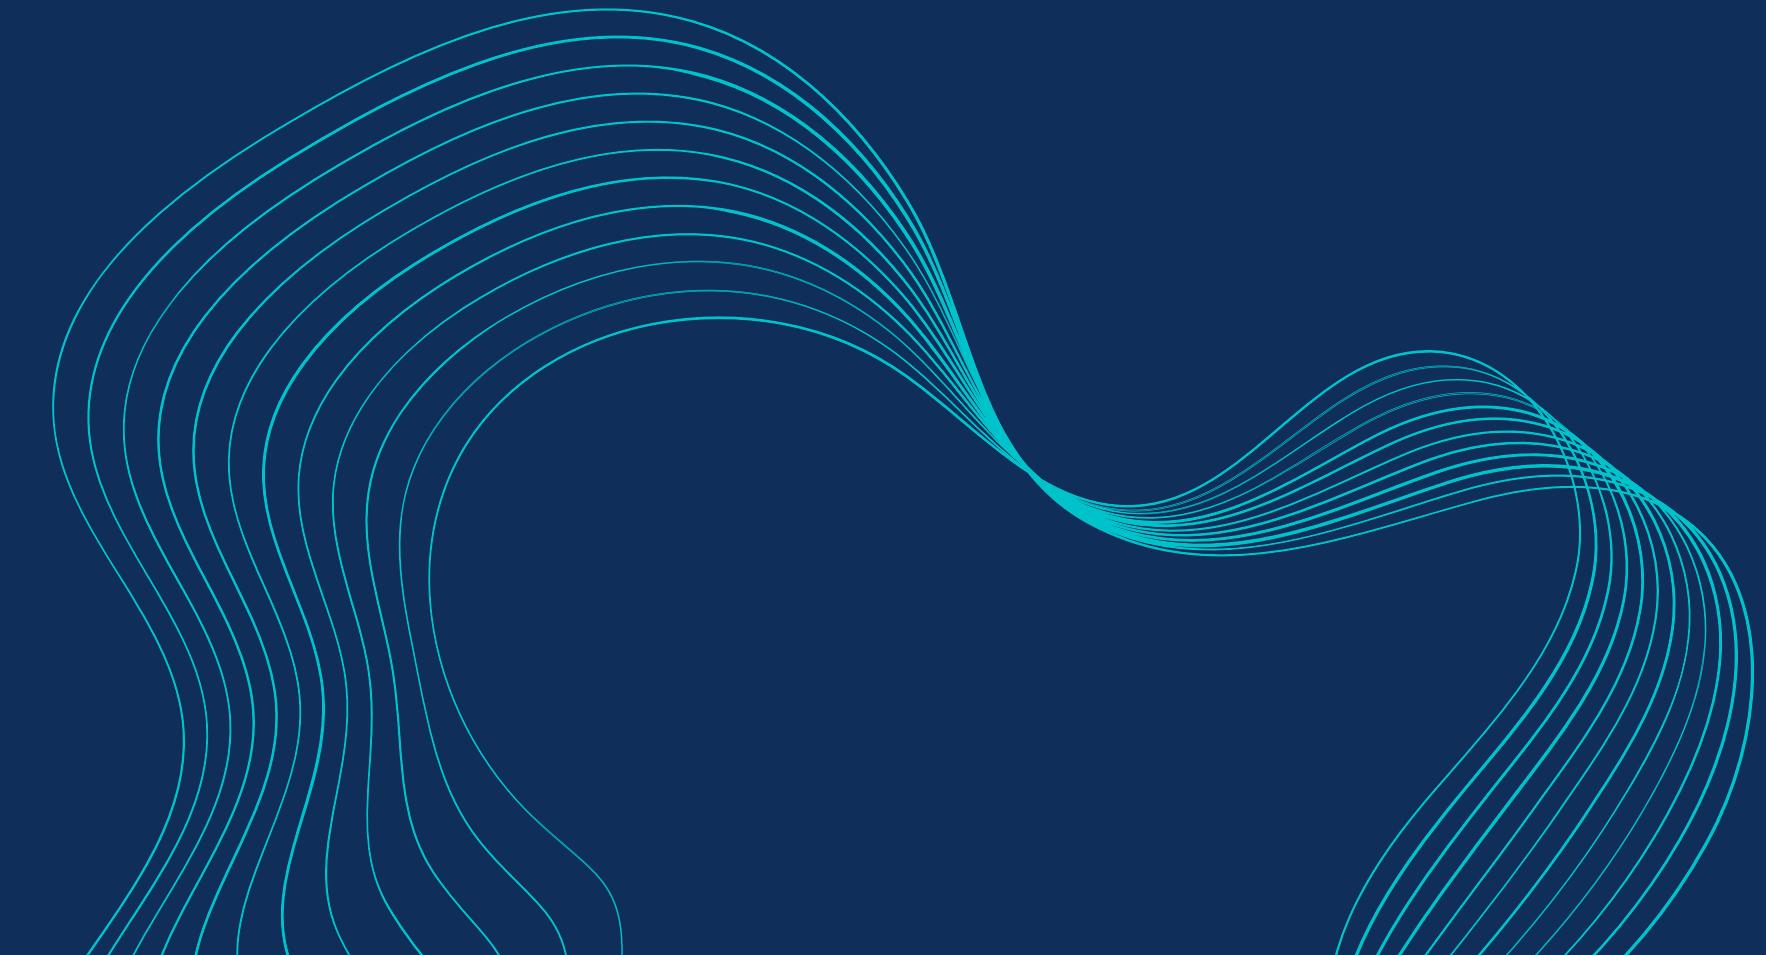
**MODEL SELECTION AND DEVELOPMENT**

4

**FINAL MODEL AND RESULTS**

5

**LIMITATIONS AND FUTURE WORK**



# Dataset

From 56 raw columns → 10 relevant, interpretable features

## Textual

- title
- descriptions

Contains seller language that directly expresses product condition (e.g. “nuevo”, “sin uso”, “usado”)

## Numerical

- price
- base\_price
- sold\_quantity
- available quantity

Capture pricing patterns and product demand, which differ between new and used items.

## Categorical

- listing\_type\_id
- buying\_mode
- category\_id

Describe how items are sold and categorized, adding contextual variation.

# Text Cleaning and Preprocessing

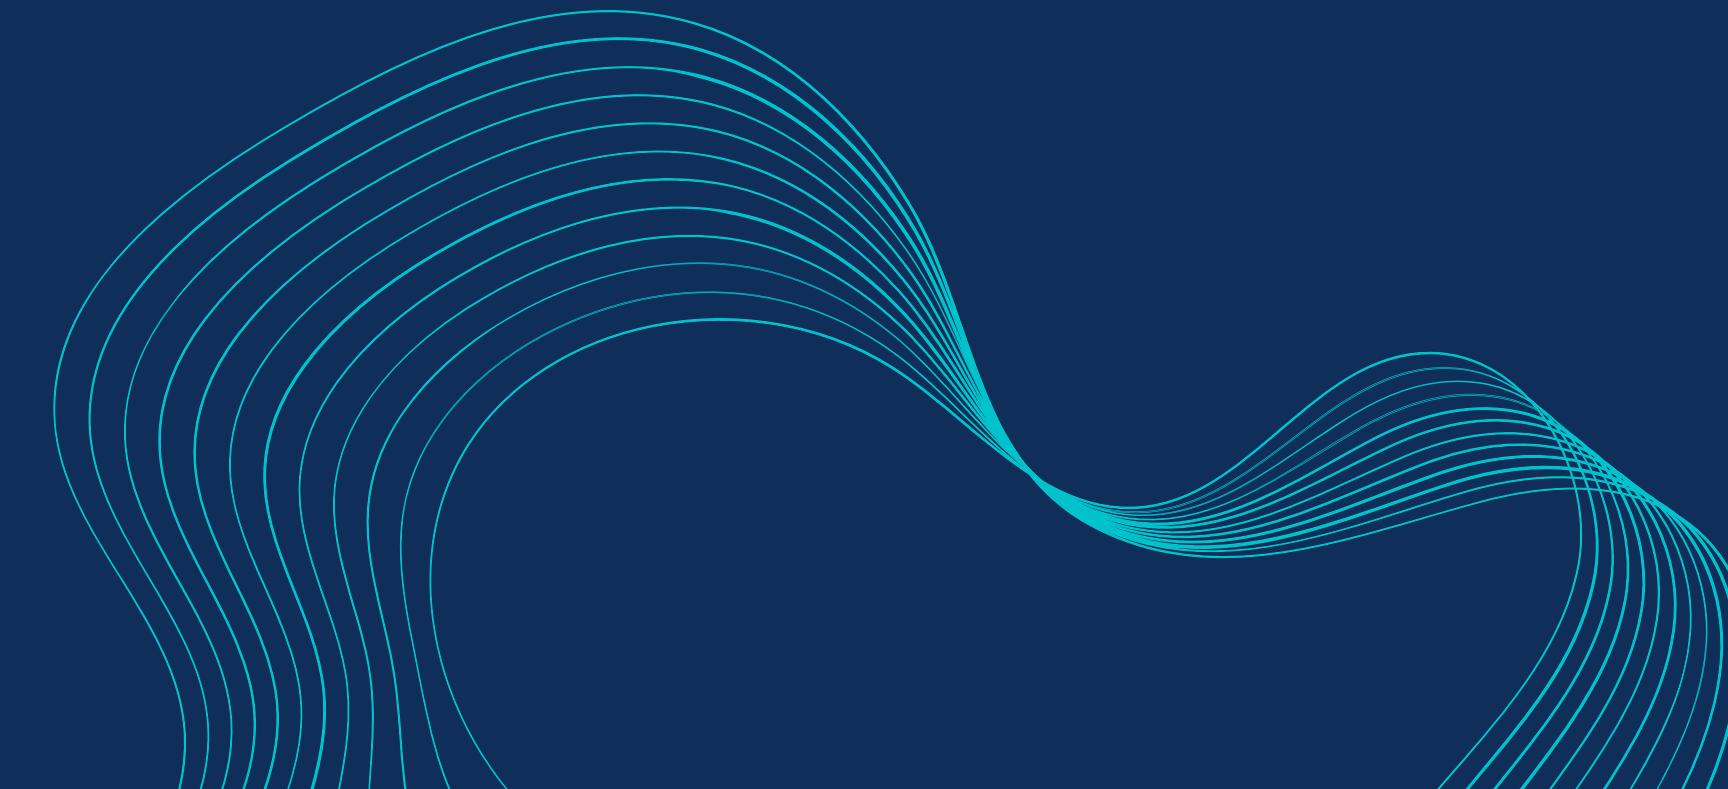
Implemented two functions:

- clean\_text()
  - lowercased text
  - removed punctuation and special characters
  - normalized accented characters
- preprocess()
  - combine titles and descriptions
  - apply the clean\_text function
  - feature engineering for the numeric and categorical variables

```
df["price_ratio"] = df["price"] / (df["base_price"] + 1)
```

```
df["log_price"] = np.log1p(df["price"].fillna(0))  
df["log_sold"] = np.log1p(df["sold_quantity"].fillna(0))
```

```
df["sold_ratio"] = df["sold_quantity"] / (df["available_quantity"] + 1)
```



# Feature groups

Three feature groups for later preprocessing techniques.

Preprocessing techniques:

- TF-IDF Vectorizer
- StandardScaler
- OneHotEncoder

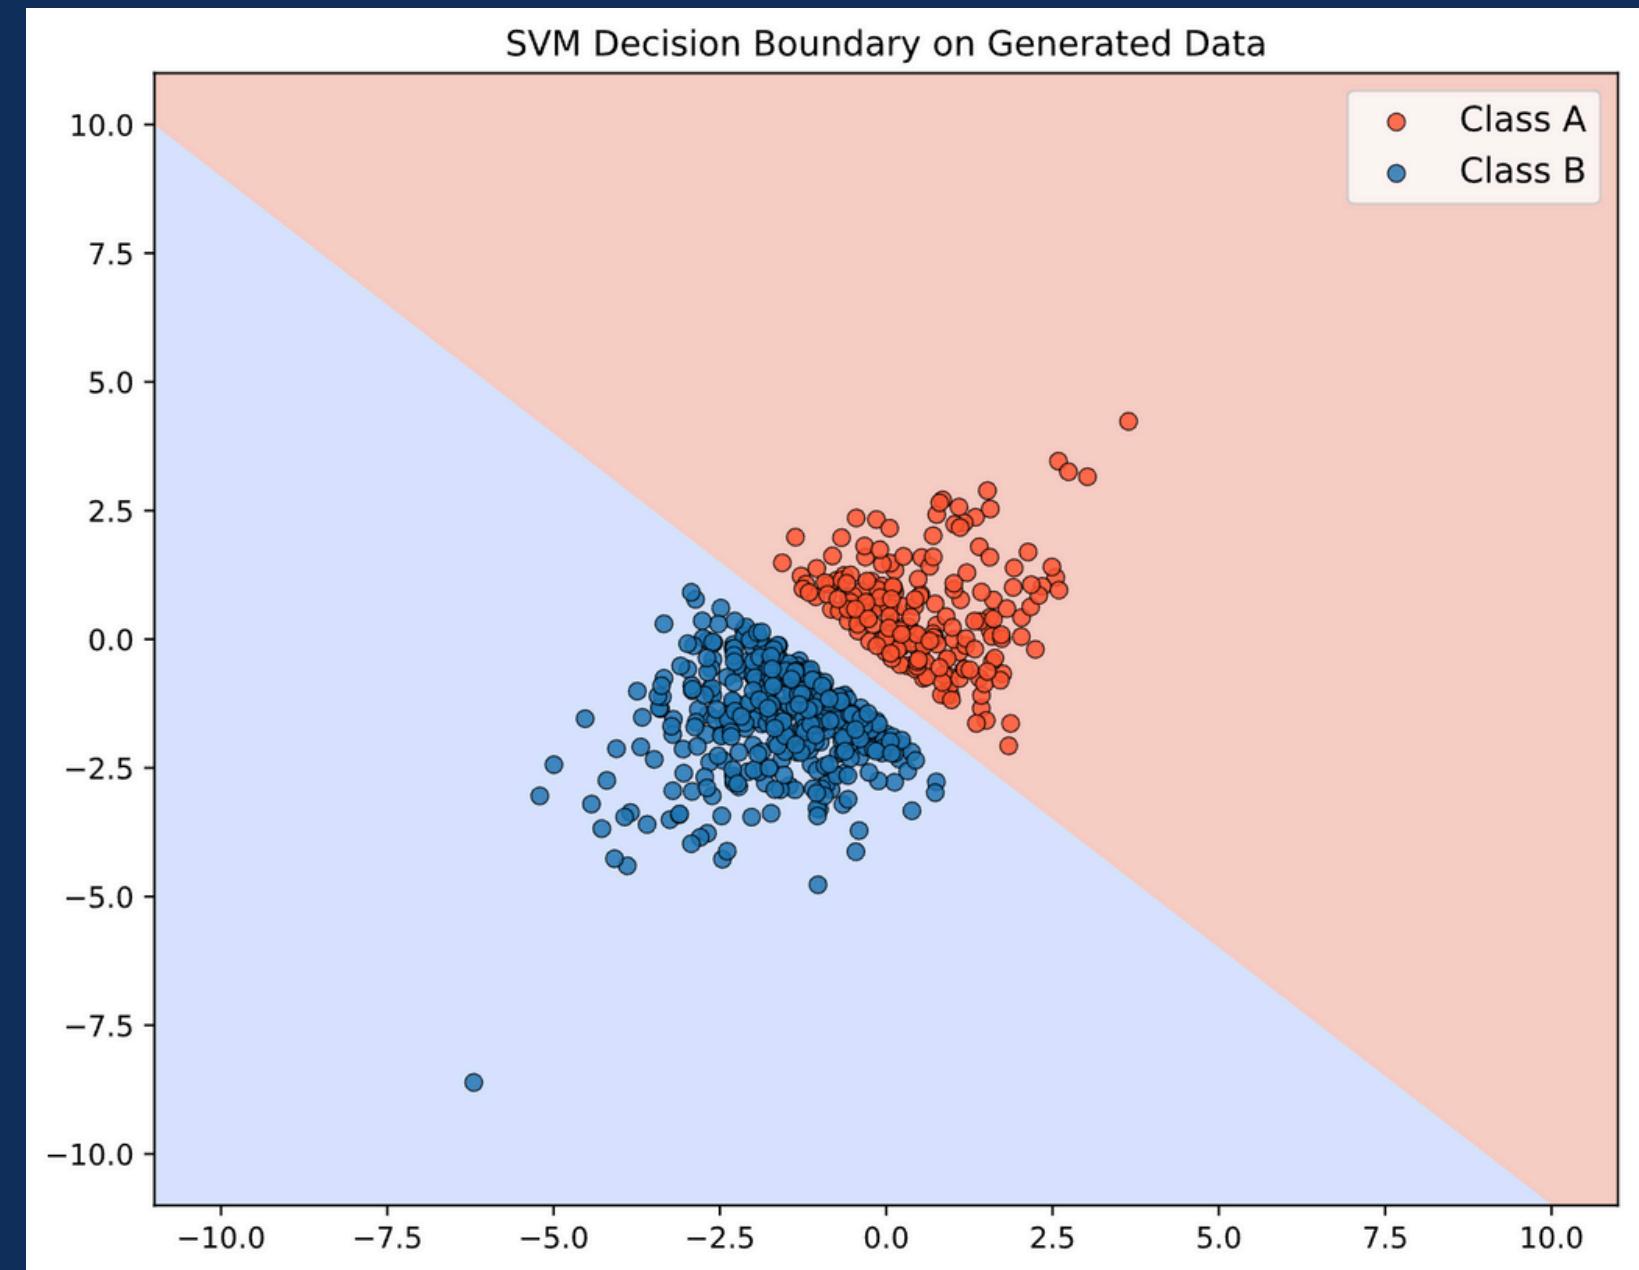
```
# Feature columns
text_col = "text"
num_cols = ["price", "base_price", "sold_quantity", "available_quantity",
            "price_ratio", "sold_ratio", "price_diff", "log_price"]
cat_cols = ["listing_type_id", "buying_mode", "category_id"]
```

# Baseline Model



# Support Vector Machines

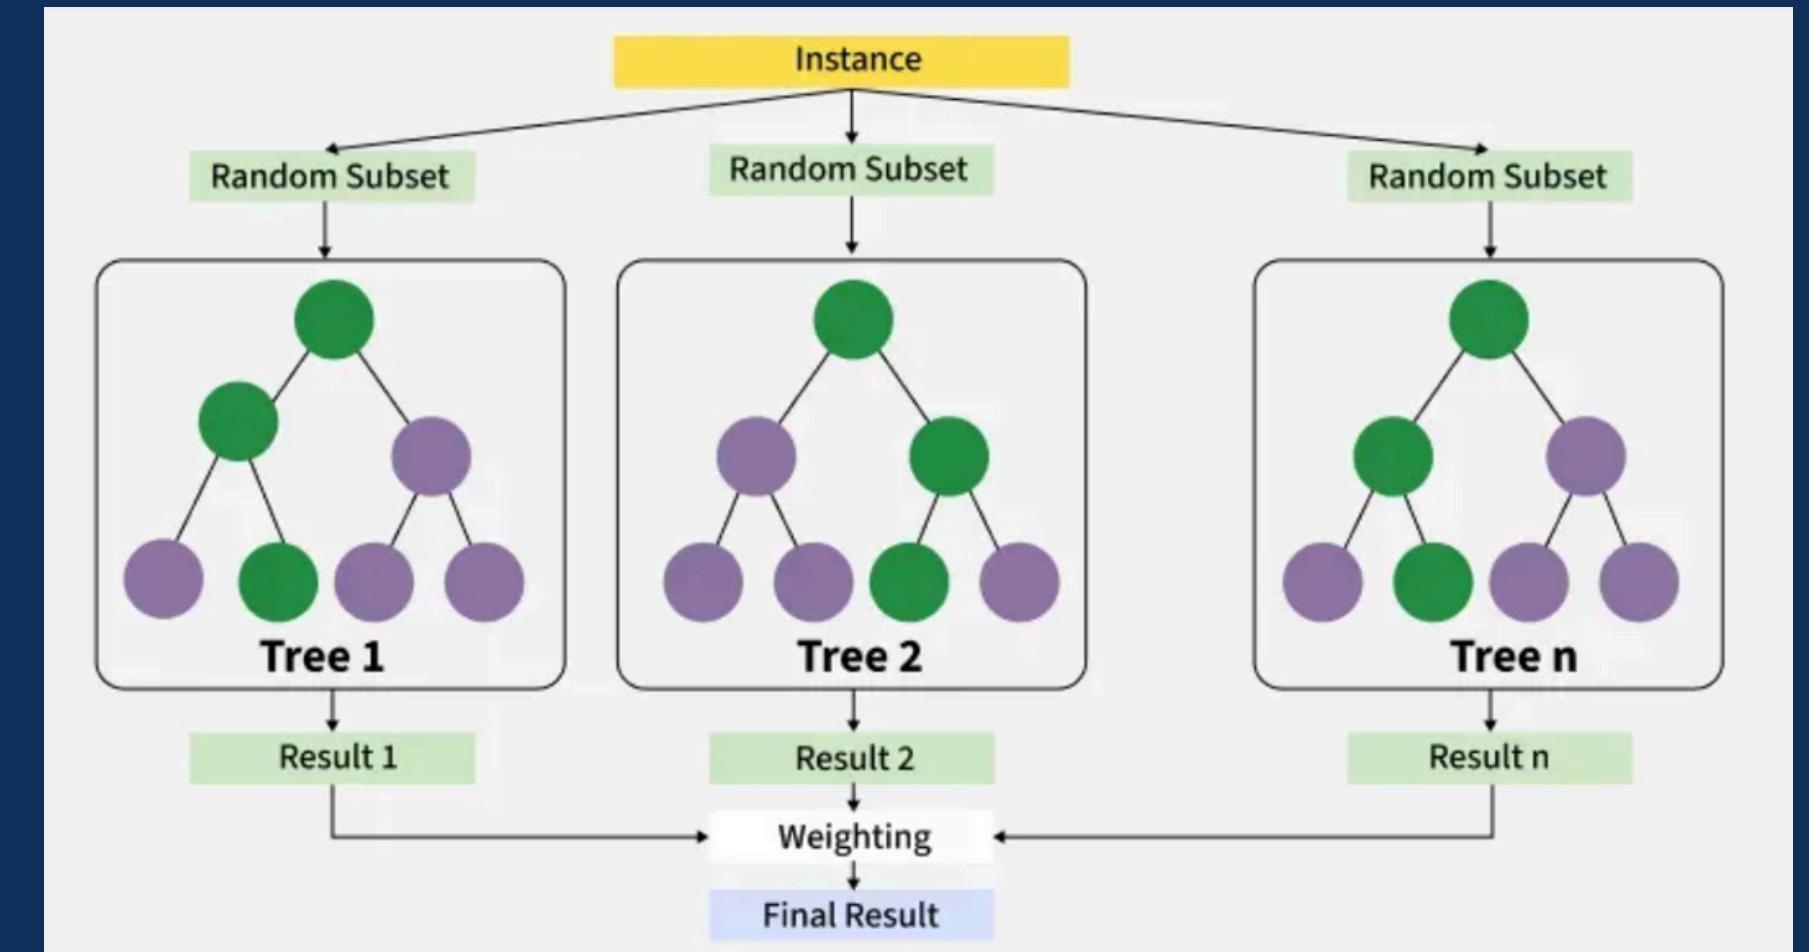
- Baseline model for comparison
- Uses TF-IDF text features
- Linear SVM: fast, interpretable, effective for high-dimensional sparse data
- Accuracy: 0.8692
- Non-linear kernel (RBF) could model complex patterns, but too heavy for this dataset.



# XGBoost

- Gradient boosting algorithm
  - builds trees sequentially
- Corrects previous errors
  - strong performance on mixed features
- Hist tree method: groups continuous features into bins
  - faster and memory-efficient
- Accuracy: 0.8828 ( $\uparrow$  from SVM)

Moving forward → improving the XGBoost



<https://www.geeksforgeeks.org/machine-learning/xgboost/>

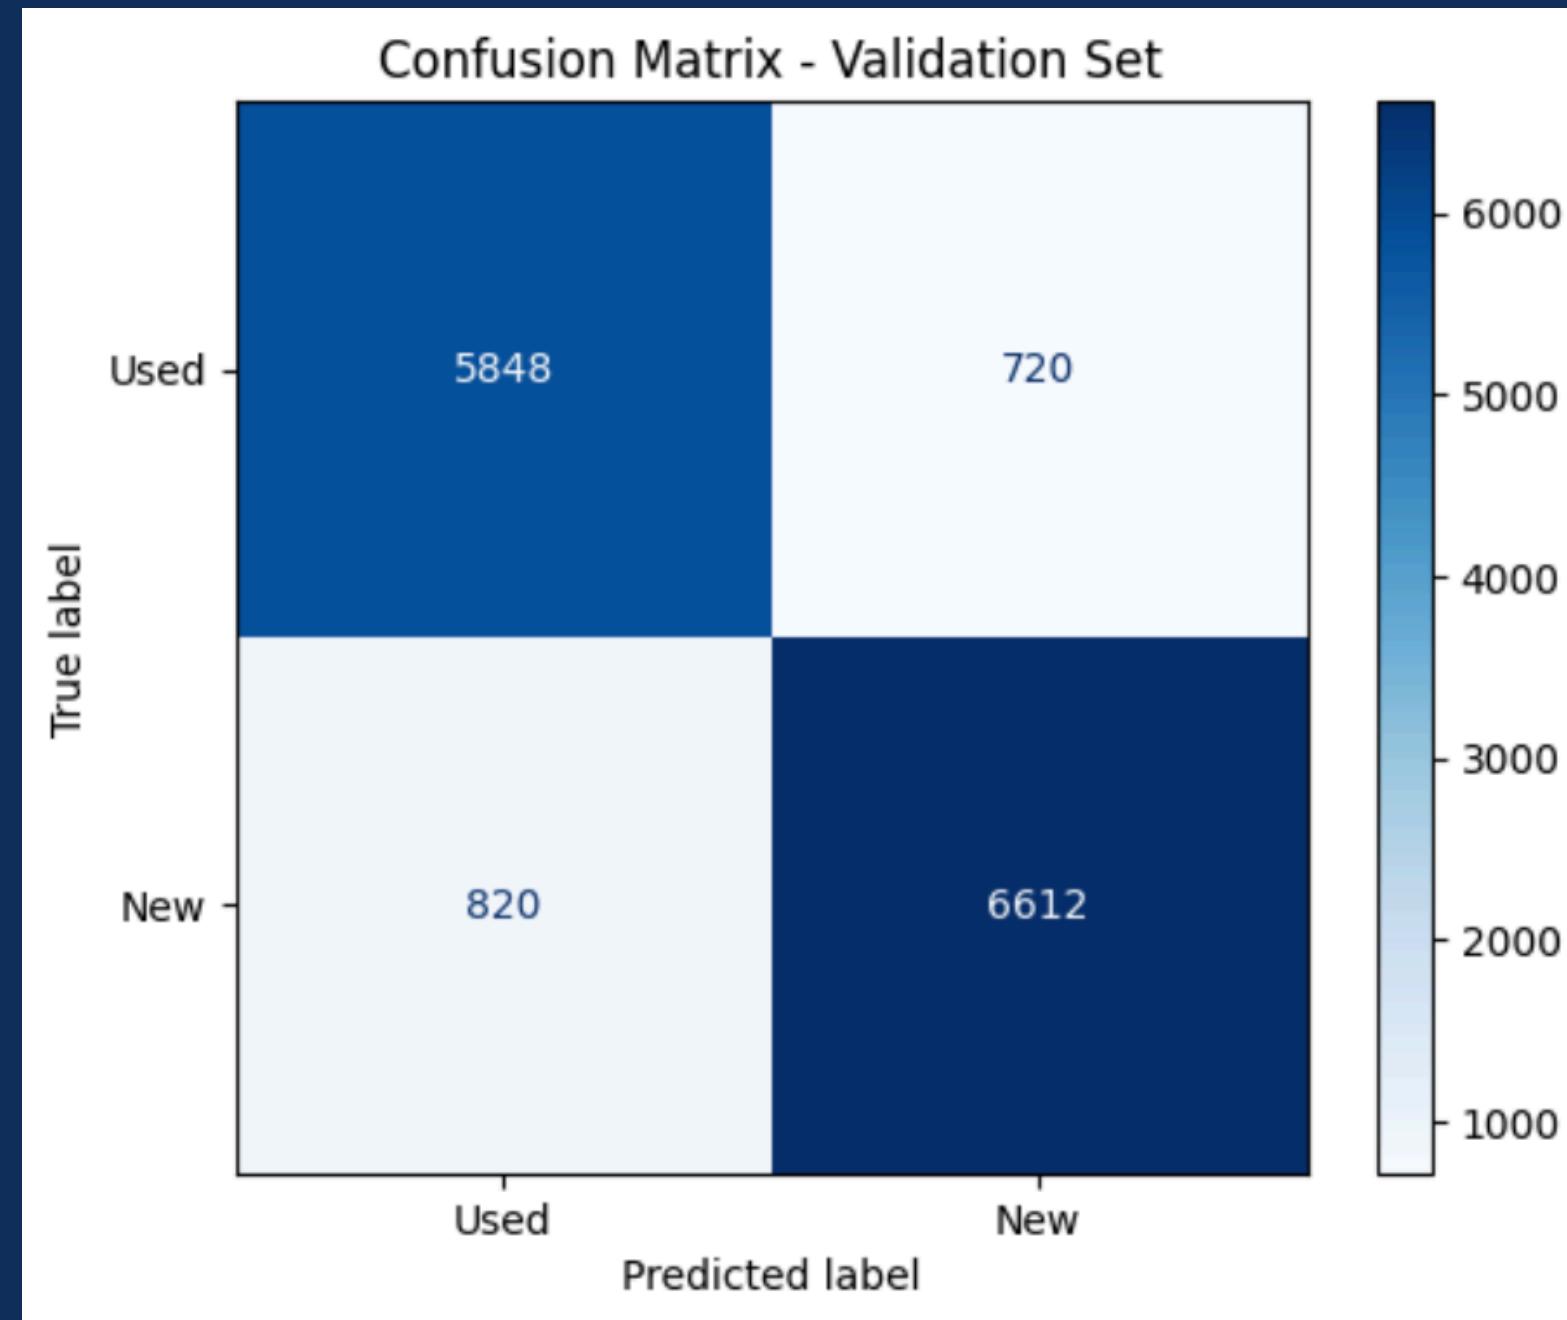
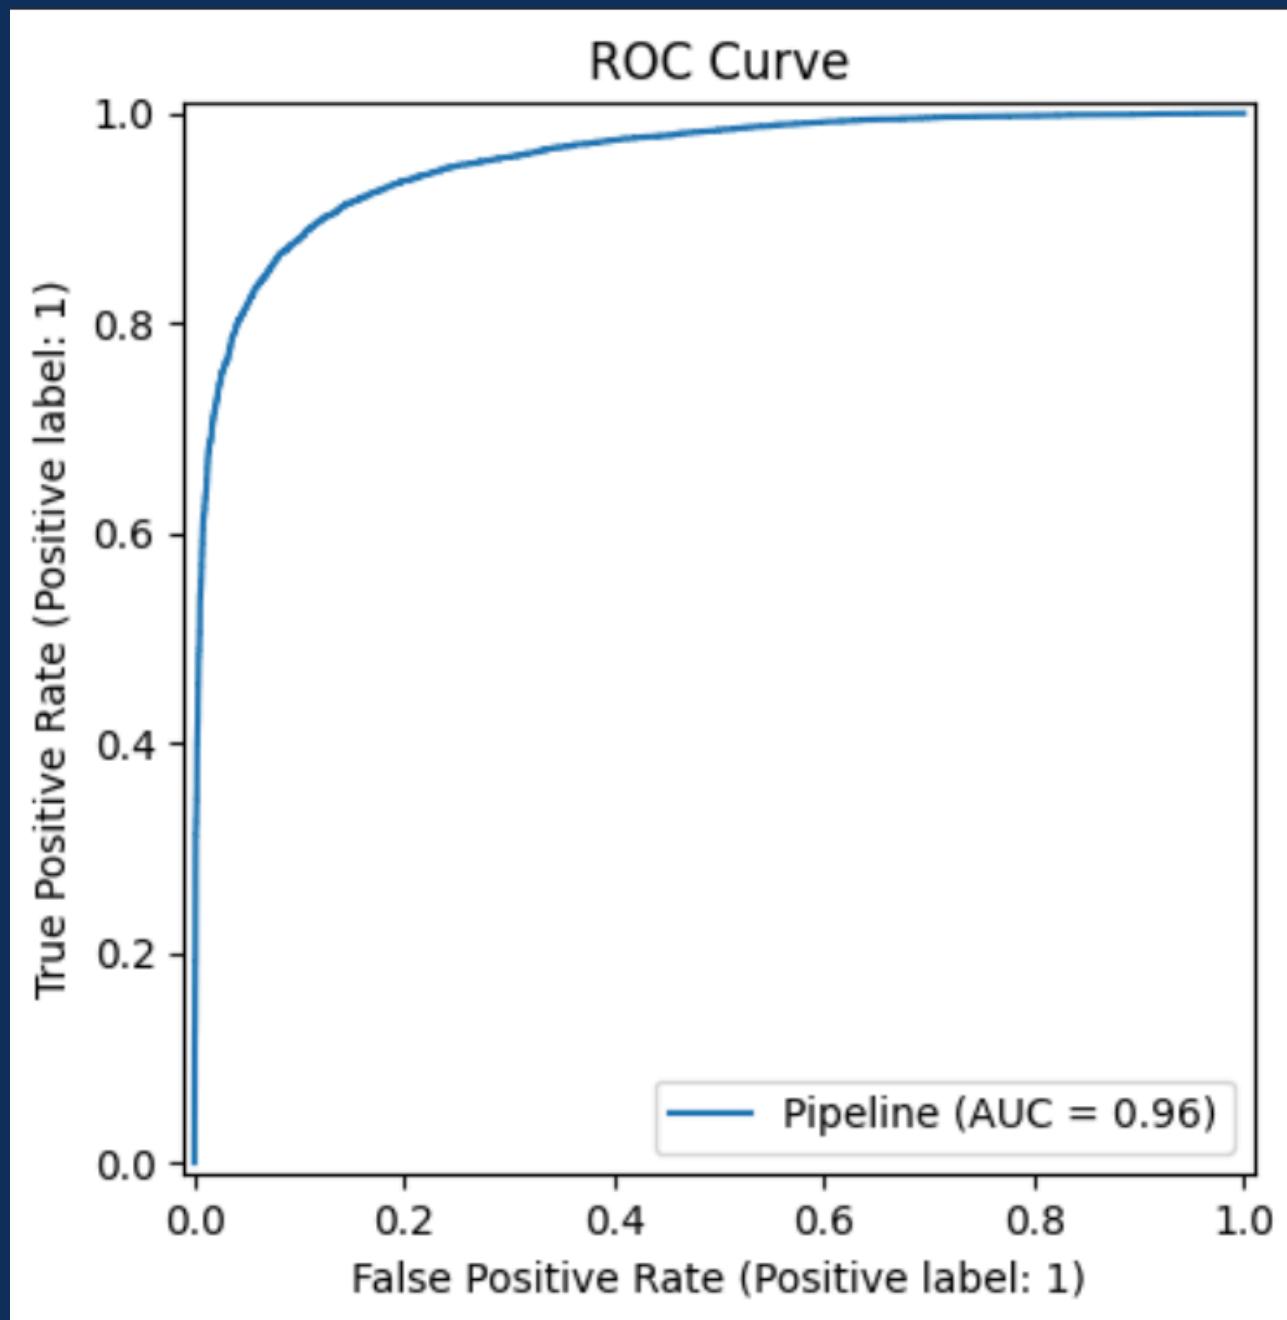
# Improvement

- ColumnTransformer → clean structure
- Added pipeline → end-to-end training
- Tuned hyperparameters
  - learning rate
  - max\_depth
  - n\_estimators

Improved model → 0.89

```
# XGBoost model
xgb_model = XGBClassifier(
    n_estimators=7000,
    learning_rate=0.03,
    max_depth=9,
    subsample=0.85,
    colsample_bytree=0.7,
    reg_lambda=1.5,
    reg_alpha=0.5,
    tree_method="hist",
    random_state=42,
    n_jobs=-1,
    eval_metric="logloss"
)
```

# Model Validation



# ■ Limitations and further work

- GridSearchCV() or Optuna for model tuning
  - Automate the search for optimal hyperparameters instead of manual tuning.
- Test additional boosting frameworks (e.g., LightGBM, CatBoost)
  - Compare performance and training efficiency on mixed feature data.
- Use richer text embeddings
  - Replace TF-IDF with word embeddings such as FastText or BERT to capture semantic meaning.
- Feature importance and interpretability
  - Analyze the most influential features to better understand what drives the model's predictions.

Gracias por su atención!